

AGIEval: A Human-Centric Benchmark for Evaluating Foundation Models

Wanjuan Zhong*, Ruixiang Cui*, Yiduo Guo, Yaobo Liang, Shuai Lu, Yanlin Wang
Amin Saied, Weizhu Chen and Nan Duan

Microsoft

{t-wzhong, v-ruicui, v-yiduoguo, yaobo.liang, shuailu}@microsoft.com
{yanlinwang, Amin.Saied, wzchen, nanduan}@microsoft.com

Abstract

Evaluating the general abilities of foundation models to tackle human-level tasks is a vital aspect of their development and application in the pursuit of Artificial General Intelligence (AGI). Traditional benchmarks, which rely on artificial datasets, may not accurately represent human-level capabilities. In this paper, we introduce AGIEval, a novel benchmark specifically designed to assess foundation model in the context of human-centric standardized exams, such as college entrance exams, law school admission tests, math competitions, and lawyer qualification tests. We evaluate several state-of-the-art foundation models, including GPT-4, ChatGPT, and Text-Davinci-003, using this benchmark. Impressively, GPT-4 surpasses average human performance on SAT, LSAT, and math competitions, attaining a 95% accuracy rate on the SAT Math test and a 92.5% accuracy on the English test of the Chinese national college entrance exam. This demonstrates the extraordinary performance of contemporary foundation models. In contrast, we also find that GPT-4 is less proficient in tasks that require complex reasoning or specific domain knowledge. Our comprehensive analyses of model capabilities (understanding, knowledge, reasoning, and calculation) reveal these models' strengths and limitations, providing valuable insights into future directions for enhancing their general capabilities. By concentrating on tasks pertinent to human cognition and decision-making, our benchmark delivers a more meaningful and robust evaluation of foundation models' performance in real-world scenarios².

AGIEval
사람을 평가하는 표준적인 시험
을 바탕으로 설계됨
- 대학 입학 시험
- 로스쿨 입학 시험
- 수학 경시대회
- 변호사 자격 시험 등

GPT4의 경우
SAT-수학, SAT, LSAT 등에서 95%
중국 국립대의 영어 시험에서 92.5%
의 성능을 보여줌
그러나 복잡한 Reasoning,
특정 도메인에 약한 모습을 보임

상황에 적절함

1 Introduction

In recent years, large foundation models, such as the latest large language models (LLMs) ChatGPT³ and GPT-4 (OpenAI, 2023), have exhibited remarkable versatility and adaptability, with plethora of applications spanning various domains as a decision-making assistant, from processing daily events to assisting in specialized fields such as law, medicine, and finance. With these advancements, AI systems are inching closer to achieving Artificial General Intelligence (AGI). As these AI systems continue to evolve and become more integrated into our daily lives, it is essential to effectively assess their general abilities in handling human-centric tasks, identify potential shortcomings, and ensure that they can handle complex, human-centric tasks effectively. Moreover, evaluating their reasoning abilities is also crucial to ensure their reliability and trustworthiness across diverse settings.

다재다능함

과잉

최근 LLM들은
사람들에게
도움이 역할을 수행중

부족하다

Traditional benchmarks for evaluating foundation models often fall short in providing an accurate assessment of their general abilities in handling human-level tasks. This is primarily due to the use

* indicates equal contribution. Yaobo Liang and Nan Duan are the corresponding authors.

²The data, code, and all model outputs are released in <https://github.com/ruixiangcui/AGIEval>
<https://chat.openai.com/chat>




Figure 1: The performance of foundation models (Text-Davinci-003, ChatGPT, and GPT-4) was evaluated on several human-centric exams under zero-shot learning with a Chain-of-Thought (CoT) prompting setting. Human performance (avg.) refers to the average performance of all test takers, while human performance (top) refers to the performance of the top 1% of test takers, except for the lawyer qualification test which uses the top 10%. Our results show that, compared to the averaged human performance, GPT-4 achieves better scores on the SAT, LSAT, and math competitions.

of artificial datasets and a lack of emphasis on real-world tasks that require human-like cognitive capabilities. Moreover, these benchmarks often focus on tasks that do not truly represent the complexities and nuances of real-world human cognition and decision-making, leading to a skewed evaluation of models' capabilities and limiting their ability to provide meaningful insights into the models' real-world applicability. Consequently, there is a growing need for a more human-centric benchmark that allows for a robust evaluation of foundation model in the context of tasks that are relevant to human reasoning and problem-solving.

In this paper, we introduce a human-centric benchmark AGIEval specifically designed to evaluate the general abilities of foundation models in tasks pertinent to human cognition and problem-solving. This benchmark is derived from official, public, and high-standard admission and qualification exams intended for general human test-takers, such as general college admission tests (e.g., Chinese College Entrance Exam (Gaokao) and American SAT), law school admission tests, math competitions, lawyer qualification tests, and national civil service exams. Furthermore, these exams are taken by a diverse range of individuals seeking entry into higher education institutions or new career paths, with millions participating annually (e.g., 12 million for the Chinese Gaokao and 1.7 million for the American SAT). As a result, these exams establish officially recognized standards for assessing human-level capabilities. Additionally, the benchmark covers bilingual tasks in both Chinese and English, allowing for a more comprehensive evaluation of the models' capabilities. By concentrating on these tasks, our benchmark provides a more meaningful and comprehensive evaluation of large language model performance in scenarios directly relevant to human decision-making. It is worth noting that since state-of-the-art foundation models, such as Text-Davinci-003, ChatGPT, and GPT-4, only have publicly available APIs for language-only tasks, we release the language-only version of AGIEval and focus on evaluating large language models in the present paper. Finally, the released AGIEval includes 20 human-centric tasks across a wide variety of subjects.

We employ 20 tasks in our newly-developed benchmark to assess the performance of cutting-edge foundation models, encompassing close-source models, i.e., GPT-4, ChatGPT and Text-Davinci-003, and an open-source model, Vicuna (Chiang et al. 2023). Our experiments explore their performance under various settings, including few-shot learning, zero-shot learning, and chain-of-thought prompting techniques. We compare the performance of these models with human performance, as illustrated in Fig. 1. Remarkably, the results reveal that GPT-4 outperforms the average human performance on LSAT, SAT, and math competitions under the zero-shot chain-of-thought (CoT) setting, demonstrating its capability on human-centric tasks. However, there remains a gap between GPT-4 and the top human performance, indicating opportunities for future improvement. We also discover that these models struggle with tasks requiring complex reasoning (e.g., LSAT-analytical reasoning and physics) or specific domain knowledge, such as law and chemistry. Moreover, our comprehensive qualitative

AGIEval은 공식적인 자격 시험을 기반으로 만들었음

- 중국 입학시험(Gaokao)
- 미국 입학시험(SAT)
- 로스쿨 입학 시험
- 수학 경시대회
- 변호사 자격시험
- 국가 공무원 시험

중국어, 미국어를 모두 커버하며 20개의 Task로 구성됨

여러 Closed-Source 모델로 평가해봄

GPT4의 경우 Zero-shot에서 인간수준의 성능을 보임

모델들은 복잡한 Reasoning (LSAT, 물리학)이나 특정 도메인(법, 화학) 등에서 약한 모습을 보인다

analyses of the ^{~을 철저하게 조사하다} four dimensions of model capabilities (i.e., *understanding, knowledge, reasoning, and calculation*) ^{다면적인} delve into their respective strengths and limitations, providing valuable insights into their general capabilities. This ^{감조하다} multi-faceted approach enables us to thoroughly examine the models' single-task behavior and identify general patterns, ultimately contributing to a more robust understanding of these state-of-the-art models and their potential applications in ^{처리하다} tackling human-level tasks.

모델의 4가지 측면을 분석하는게 목적임

- Understanding
- Knowledge
- Reasoning
- Calculation

AGI를 위해 모델의 어떤 점이 부족한지 평가해주는게 목적

We aim to drive innovation in developing more effective and reliable AI assistants that advance towards Artificial General Intelligence (AGI). By identifying areas for improvement and understanding their limitations, we can enhance the models' performance and gain a deeper understanding of their underlying mechanisms. Ultimately, this will lead to more powerful and dependable AI systems that better serve the needs of users across a wide range of applications.

In summary, our research ^{감조하다} underscores the importance of evaluating foundation models in the context of human-level tasks and provides a solid benchmark for such evaluations. We hope our findings inspire further innovation and progress in the development of large foundation models, ultimately leading to more reliable and effective AI systems.

2 Background and Related Work

2.1 Large Foundation Model

In recent years, large foundation models, like Large Language Models (LLMs) (e.g., GPT-3 (Brown et al. 2020), GPT-4 (OpenAI, 2023), OPT (Zhang et al., 2022a) and FLAN-T5 (Chung et al., 2022)) have successfully demonstrated unprecedented performance in a wide range of natural language tasks. The success of these models can be attributed to advances in deep learning techniques, architectural improvements, and the availability of massive amounts of data for training. The most recent cutting-edge language models, such as GPT-4 (OpenAI 2023) and ChatGPT⁴, have continued to demonstrate substantial adaptability to a diverse array of tasks and domains and have served as a daily decision-making assistant for human beings.

However, despite their impressive performance on various benchmarks, concerns have been raised about the reasoning abilities, trustfulness and real-world applicability of these models (Marcus & Davis 2019). Therefore, there is still a need for more comprehensive benchmarks that assess the general reasoning abilities of large foundation models in human-centric tasks.

2.2 Existing Benchmarks

Constructing benchmarks is a reliable way to establish evaluation standards and monitor model performance. Numerous benchmarks (Thorne et al., 2018; Rajpurkar et al., 2016; Bowman et al., 2015; Ehrmann et al., 2022) have been proposed and widely adopted for evaluating single-task performance, such as SQuAD (Rajpurkar et al., 2016) for assessing answer extraction ability and SNLI (Bowman et al., 2015) for evaluating natural language inference capability.

SQuAD
답변 추출 능력 평가

SNLI
자연어 추론 능력 측정

The emergence of general language models (LMs) like BERT (Devlin et al., 2019) has made it increasingly essential to develop more comprehensive benchmarks to assess the general capabilities of these LMs. GLUE (Wang et al., 2018) and SuperGLUE (Wang et al., 2019) are popular benchmarks that evaluate language model performance across diverse NLP tasks, including paraphrase identification, sentiment analysis, and more. GLUE series benchmarks have significantly influenced language model development, encouraging researchers to enhance their models' generalization capabilities. The LAMBADA language modeling task (Paperno et al., 2016) assesses language models' ability to capture long-range dependencies in text. SentEval (Conneau & Kiela, 2018) and DecaNLP (McCann et al., 2018) also set benchmarks for evaluating models' general capabilities.

GLUE, SuperGLUE
여러 NLP Task 평가
- Paraphrase Identification
- Sentiment Analysis
등등

LAMBADA
텍스트에서 긴 범위의 의존성을 파악할 수 있는지 평가

SentEval / Deca NLP
모델의 일반적 능력 평가

Despite their broad applicability, these benchmarks mainly consist of artificially curated datasets designed to evaluate specific machine skills, rather than real-world problems aimed at assessing human behaviors. Consequently, these benchmarks primarily focus on simpler textual understanding rather than complex reasoning abilities aligned with real-world applicability.

⁴<https://chat.openai.com/chat>

MMLU (Hendrycks et al., 2020) addresses this issue by collecting questions from online sources covering a diverse set of subjects (e.g., history, humanities) that humans learn, pushing towards human-centric evaluation. However, there are key differences between MMLU and our work: (1) **Source of benchmark:** Our benchmark is derived from high-standard and official human-centric exams, such as college admission tests and professional qualification tests, ensuring a robust and well-standardized evaluation of language models. In contrast, MMLU’s data source is not explicitly mentioned, and it is unclear whether the tasks come from similarly professional and high-quality sources. (2) **Bilingual Evaluation:** Our benchmark is bilingual, containing both English and Chinese tasks, which allows for a more comprehensive assessment of language models across different languages and cultures. In contrast, MMLU only contains English data, limiting its evaluation scope to a single language.

MMLU
Human-Centric 평가

AGIEval이랑 다른점:

1) AGIEval은 수준높은 평가 시험에 기반함
MMLU는 Source가 명시되어 있지 않음

2) 영어 & 중국어 모두 가능

2.3 Evaluation of Recent Large Foundation Models

The rapid development of foundation models has significantly propelled research efforts in evaluating their performance and behavior. As these models continue to grow in size and complexity, it becomes increasingly important to assess their abilities in understanding textual data and performing complex reasoning and problem-solving tasks that humans typically excel at. Consequently, studies have been conducted to understand their strengths and limitations.

최근 LLM들은
Complex Reasoning을 평가하는게 관심임

ToxiGen (Hartvigsen et al., 2022) and BOLD (Dhamala et al., 2021) evaluates the bias in language models. Liang et al. (2022) conducted a holistic evaluation of dozens of language models (up to InstructGPT (Ouyang et al., 2022)) on several previous NLP tasks (e.g., Openbook QA (Kwiatkowski et al., 2019), Mihaylov et al., 2018) and text classification (Clark et al., 2019)), focusing on model behavior analysis. However, this work does not cover the most recent LLMs such as ChatGPT and GPT-4, and evaluations are mainly conducted on artificially curated NLP datasets rather than real-world scenarios for testing humans.

ToxiGen, BOLD 등은 모델의 Bias를 평가함

Recognizing this limitation, recent reports have emphasized the importance of evaluations in human-centric scenarios. For example, Choi et al. (2023) assessed ChatGPT’s performance using essay questions from law schools. Recent reports start to investigate the capabilities of GPT-4, like Bubeck et al. (2023) and OpenAI (2023). The official technical report of GPT-4 (OpenAI, 2023) also underscored the importance of evaluating models’ behaviors on human exams and analyzed GPT-4’s performance on several such exams. However, the relevant benchmarks in these reports and the corresponding model outputs are not publicly available, and the evaluation metric is also not transparent. These factors limit further research to follow up their evaluation and to make direct comparison.

최근의 리포트들은
Human-Centric Scenario에서 평가하는 것을 중요시함

GPT4 테크니컬 리포트에서도 Human exam의 퍼포먼스를 평가하기도 함

To bridge this gap, we propose AGIEval, which releases collected questions from high-standard, official human exams and supports a standardized, automatic evaluation metric. Furthermore, we conduct comprehensive experiments with the most recent LLMs including ChatGPT and GPT-4 and release the model outputs to foster analysis and research within the entire community.

3 Human-Centric Benchmark

In this section, we outline the design principles and process of constructing our human-centric benchmark.

3.1 Design Principles

Emphasis on human-level cognitive tasks: In designing the human-centric benchmark, our primary objective is to center on tasks that closely align with human cognition and problem-solving, ultimately assessing the general abilities of foundation models in a more meaningful and comprehensive manner. To achieve this, we incorporate a diverse range of official, public, and high-standard admission and qualification exams that cater to general human test-takers. These exams, which include college admission tests, law school admission tests, math exams, lawyer qualification tests, and national civil service exams, are taken by millions of individuals seeking entry into higher education or new career paths. By adhering to these officially recognized standards for evaluating human-level capabilities, our benchmark design ensures that the assessment of model performance is directly relevant to human decision-making and cognitive abilities.

Human-Centric Benchmark

- 대학 입학시험
- 로스쿨 입학시험
- 수학 시험
- 변호사 자격 시험
- 공무원 시험

등 사람이 실제로 보는 시험을 기반으로 벤치마크 작성

Relevance to real-world scenarios: The second design principle guiding the construction of our human-centric benchmark is the emphasis on tasks that hold significant implications for real-world situations. By selecting tasks derived from high-standard admission and qualification exams, we ensure that the assessment reflects the complexity and practicality of challenges that individuals commonly encounter in various fields and contexts. This approach allows us to not only gauge the models' performance in relation to human cognitive abilities but also better understand their applicability and effectiveness in real-life scenarios. Consequently, our benchmark design contributes to the development of AI systems that are more reliable, practical, and well-suited for addressing a wide range of real-world problems.

사람이 실제로 보는
평가 시험을 기반으로
하기 때문에
측정하다
Real-Life에
도움이 되는지 측정하는
Benchmark의 역할을
수행할 수 있음

3.2 Exam Selection

In line with the design principles outlined above, we have selected a diverse range of standardized, high-quality exams that emphasize human-level reasoning and real-world relevance. Some exams are participated by millions of human test-takers annually. For example, 12 millions of students participate in Gaokao every year. Statistics of annual human participants are reported in Tab. 1. The following categories of human-centric exams are included in our benchmark:

General College Entrance Exams: These exams, include Graduate Record Examinations (GRE), Scholastic Assessment Test (SAT) and China College Entrance Exam (Gaokao), are designed to assess the general aptitude and subject-specific knowledge of students seeking entry into higher education institutions. College admission tests encompass various subjects and require critical thinking, problem-solving, and analytical skills, making them ideal for evaluating the performance of large language models in relation to human cognition. More specifically, we collect exams corresponding to 8 subjects from Chinese Gaokao: history, math, English, Chinese, geography, biology, chemistry and physics. We select mathematical questions from GRE, select English and math subjects from SAT to construct the benchmark.

Exam 고르기
1) 대학 입학시험
GRE: 수학
SAT: 영어, 수학
Gaokao: 8개 과목
(역사, 수학, 영어, 중국어,
지리학, 생물학, 화학, 물리학)

Law School Admission Test: Law school admission tests, such as the LSAT, are intended to measure the reasoning and analytical skills of prospective law students. These tests include sections on logical reasoning, reading comprehension, and analytical reasoning, which challenge the test-takers' ability to analyze complex information and draw accurate conclusions. Incorporating these tasks in our benchmark enables us to assess language models' capabilities in legal reasoning and analysis.

2) 로스쿨 입학시험
LSAT
논리적 추론, 독해, 분석 추론 등
복잡한 상황을 분석하여
정확한 결론을 내리는
능력을 테스트할 수 있음

Lawyer Qualification Test: Lawyer qualification tests, such as the bar exam, assess the legal knowledge, analytical skills, and ethical understanding of individuals pursuing a career in law. These exams cover a broad range of legal topics, including constitutional law, contract law, criminal law, and property law, and require candidates to demonstrate their ability to apply legal principles and reason effectively. By incorporating lawyer qualification tests in our benchmark, we can evaluate language models' performance in the context of professional legal expertise and ethical judgment. Specifically, we select questions from previous Chinese lawyer qualification tests to build the benchmark.

3) 변호사 자격시험
bar exam
법률적 지식, 윤리적 이해를
테스트할 수 있음

Graduate Management Admission Test (GMAT): The GMAT is a standardized exam designed to assess the analytical, quantitative, verbal, and integrated reasoning skills of prospective graduate business school students. The GMAT consists of sections such as Analytical Writing Assessment, Integrated Reasoning, Quantitative Reasoning, and Verbal Reasoning, which evaluate the test-takers' abilities to think critically, analyze data, and communicate effectively. Incorporating GMAT tasks in our benchmark enables us to assess the language models' performance in a business context and their potential to assist in decision-making and problem-solving in management scenarios.

4) 경영대학원 입학시험
GMAT
여러가지 추론 능력,
분석적 능력을 테스트함

High School Math Competitions: High school math competitions, such as the American Mathematics Competitions (AMC) and the American Invitational Mathematics Examination (AIME), challenge students' mathematical abilities, creativity, and problem-solving skills. These contests cover a wide range of mathematical topics, including number theory, algebra, geometry, and combinatorics, and often feature unconventional problems that require inventive approaches to solve. By incorporating tasks from high school math competitions into our benchmark, we can further evaluate the language models' aptitude for tackling complex and creative mathematical problems, as well as their ability to generate novel solutions.

5) 고등학교 수학 경시 대회
AMC
AIME
다양한 수학 과목 평가
(대수학, 선형대수학, 기하학...)

Chinese Civil Service Examination: The Chinese Civil Service Examination is a standardized test administered in China to assess the aptitude and skills of individuals seeking entry into the civil service. These exams evaluate a range of competencies, such as general knowledge, reasoning

6) 중국 공무원 시험
Chinese Civil Service exam
일반 상식, 추론 능력,
언어 능력 및
본인 직무에 맞는 전문 지식 평가

abilities, language skills, and subject-specific expertise related to the roles and responsibilities of various civil service positions in China. Including tasks from national civil service exams in our benchmark allows us to gauge the language models’ performance in public administration contexts and their potential to contribute to policy-making, decision-making, and public service delivery processes.

By including these tasks in our human-centric benchmark, we create a comprehensive assessment that effectively measures the performance of foundation models in various real-world, human-level reasoning scenarios.

3.3 Construction of Human-Centric Benchmark

3.3.1 Data Collection

As previously mentioned, our human-centric benchmark comprises questions from a diverse range of official and high-quality exams, originally designed for human test-takers. These exams include general college admission tests (GRE, Gaokao, SAT), entrance exams for specific majors (such as LSAT and GMAT), high school math competitions (AMC and AIME), as well as the national civil service examination and lawyer qualification test in China.

주관적인 답변을 배제하기 위해
객관적인 답변셋만 사용함

- 다지선다형 문제
- 빈칸 채우기 문제

Since evaluating model performance on subjective questions is challenging without human expert scoring, we believe such questions are unsuitable for inclusion in this benchmark for consistent assessment. To ensure a robust and standardized evaluation metric, we have removed all subjective questions, retaining only objective ones, such as multiple-choice and fill-in-the-blank questions.

With regard to data collection, we gather Gaokao and SAT questions from publicly available online sources, along with their corresponding solutions or explanations. For the LSAT, we utilize data from Wang et al. (2022) and Zhong et al. (2022), which encompasses three tasks (logical reasoning, reading comprehension, and analytical reasoning) from the LSAT administered between 1991 and 2016. For Chinese civil service examinations, we repurpose data from LogiQA (Liu et al. 2021), a dataset built on various types of logical reasoning questions collected from the National Civil Servants Examination of China. It is worth noting that LogiQA consists of bilingual questions (English and Chinese), where the English version is a translated version of the original Chinese version.

Gaokao, SAT
웹에서 공개된 시험자료와
해설지를 받아서 구축

LSAT
3개의 Task
(논리 추론, 독해, 분석추론)
1991년 ~ 2016년 데이터 수집

중국 공무원시험
LogiQA 데이터셋 사용
영어 / 중국 버전이 있음

For high school math competitions, we employ data from the MATH dataset (Hendrycks et al.), comprising questions from AMC and AIME. Furthermore, we incorporate GRE and GMAT questions from AQuA-RAT (Ling et al. 2017), which emphasizes algebraic word problems. In the case of the Chinese Civil Service Examination, we reuse instances from JEC-QA (Zhong et al. 2020), a large-scale dataset derived from the National Judicial Examination of China. We down-sample the two types of JEC-QA and MATH to 1,000 instances each.

고등학교 수학경시대회
MATH 데이터셋 사용
(AMC, AIME 포함됨)

GRE, GMAT
AQuA-RAT에서 사용

중국 공무원 시험중
중국 판사 시험 모음인
JEC-QA 데이터셋 사용
총 8062개 질문셋

As a result, we construct a benchmark consisting of 8,062 questions for evaluation. Detailed data statistics are presented in Table I. It is worth noting that our benchmark is bilingual, encompassing both English and Chinese tests. This design enables the evaluation of a broader scope of model capabilities, reflecting their performance and adaptability across different languages.

3.3.2 Evaluation Metrics

The benchmark questions consist of objective formats: multiple-choice and fill-in-the-blank questions. For multiple-choice questions, we adopt standard classification accuracy as the evaluation metric. For fill-in-the-blank questions, we employ Exact Match (EM) and F1 metrics.

Evaluation Metric으로는

- 다지선다형: Accuracy
- 빈칸채우기: EM, F1

3.3.3 Human Performance

It is worth noting that we also report the average and top performances of human test-takers as human-level boundaries for each task. The human performance results for MATH, LogiQA, and JEC-QA were gathered from their respective publications, while for AQuA-RAT (GMAT, GRE), LSAT, SAT, and Gaokao, we estimated the human performances by scaling the scores of the average (50%) and top (1%) test takers against the full scores to 100. It is important to note that while our datasets and human accuracies provide a useful approximation of human test-taker ability, they do not fully represent the range of skills and knowledge that such individuals may possess.

시험문제이기 때문에
각 Task마다 평균점수로
Human-level boundary
설정이 가능함

평균이기 때문에
개개인의 능력을
평가한 것이 아닌점 유의

Table 1: Introduction of the exams included in AGIEval. We highlight the number of human participants taking these exams annually (column “# Participants”). We also report the number of instances and average token number in AGIEval.

Exams	#Participants	Language	Tasks	Subject	# Instance	#Avg. Token
Gaokao	12M	Chinese	GK-geography	Geography	199	144
			GK-biology	Biology	210	141
			GK-history	History	243	116
			GK-chemistry	Chemistry	207	113
			GK-physics	Physics	200	124
			GK-En	English	306	356
			GK-Ch	Chinese	246	935
			GK-Math-QA	Math	351	68
SAT	1.7M	English	GK-Math-Cloze	Math	118	60
			SAT-En.	English	206	656
Lawyer Qualification Test	820K	Chinese	SAT-Math	Math	220	54
			JEC-QA-KD	Law	1000	146
Law School Admission Test (LSAT)	170K	English	JEC-QA-CA	Law	1000	213
			LSAT-AR	Law-Analytics	230	154
			LSAT-LR	Law-Logic	510	178
Civil Service Examination	2M	English	LSAT-RC	Law-Reading	260	581
			LogiQA-en	Logic	651	144
GRE	340K	English	LogiQA-ch	Logic	651	242
GMAT	150K	English	AQuA-RAT	Math	254	77
AMC	300K	English	MATH	Math	1000	40
AIME	3000	English				

4 Evaluation of Foundation Models

4.1 Foundation Models Selected for Evaluation

In this section, we evaluate the performance of various state-of-the-art language models on our benchmark dataset. The selected models represent a diverse range of capabilities and are widely recognized in the field. The language models chosen for evaluation include:

- **GPT-4**: The fourth-generation model in the GPT series, GPT-4 is a state-of-the-art, large-scale generative pre-trained transformer model with enhanced performance and a broader knowledge base compared to its predecessors. While it does not surpass human capabilities in many real-world scenarios, it demonstrates human-level performance on many scenarios. GPT-4 has been iteratively aligned using insights from adversarial testing and ChatGPT, resulting in remarkable improvements in factuality,
- **ChatGPT**: A conversational AI model developed by OpenAI, ChatGPT is designed to engage in interactive and dynamic conversations. It has been trained on a vast instruction dataset and further tuned by reinforcement learning with human feedbacks (RLHF), enabling it to provide contextually relevant and coherent responses aligned with human expectation.
- **Text-Davinci-003**: GPT-3.5 is an intermediate version between GPT-3 and GPT-4, offering improved performance over GPT-3 through further instruction tuning. It serves as a bridge between the two models, allowing for comparative analysis. Specifically, we select the variant **Text-Davinci-003** from GPT-3.5 series for evaluation.
- **Vicuna-13B (Chiang et al., 2023)**: It is an open-source LLM, trained on user-shared conversations from ShareGPT by fine-tuning LLaMA. It achieves over 90% of the quality of OpenAI’s ChatGPT. As of May 2023, Vicuna-13B ranks among the top models on the Open LLM Leaderboard (Beeching et al., 2023).

A recent technical report (SenseTime, 2023) further evaluates foundation models (GLM-130B (Du et al., 2022), LLaMa-65B (Touvron et al., 2023) and InternLM-104B (SenseTime, 2023)) on AGIEval.

FM Evaluation 해설

- GPT4

- ChatGPT : RLHF한 모델

- Text-Davinci-003: GPT3.5 모델

- Vicuna-13B: LLaMA에 ShareGPT의 유저 대화 데이터를 학습

그 외에도 GLM-130B, LLaMA-65B, InternLM-104B도 평가

In the following sections, we will discuss the evaluation methodology, including the metrics used to assess the performance of these models, and present a comparative analysis of their performance on the benchmark dataset.

4.2 Experimental Setup

In this section, we describe the experimental setup used to evaluate the performance of large language models, including GPT-4, ChatGPT and Text-Davinci-003, on our human-centric reasoning tasks benchmark. To gauge the adaptability of the language models, we conducted two types of evaluations: zero-shot and few-shot. We further implement a “Chain-of-Thought” reasoning evaluation for both zero-shot and few-shot learning. Fig. 2 describes the concrete prompting examples for zero-shot testing, few-shot testing and chain-of-thought prompting. To foster analyses about models’ behavior and foster relevant research, we release the model outputs under all the settings.

실험은 4가지 프롬프트로 진행함

- Zero-shot
- Zero-shot + CoT
- Few-shot
- Few-shot + CoT

Zero-shot은 모델의 선천적인 능력을 평가함

Few-shot은 모델이 주어진 정보를 금방 습득하고 적용하는 능력을 평가함

4.2.1 Zero-shot and Few-shot Evaluation

In the zero-shot setting, models were directly evaluated on the questions without being provided any prior examples of the specific tasks. This scenario tests the models’ innate ability to reason and solve problems without explicit training.

선천적인

In the few-shot setting, models were given a small number of examples (e.g., 5) from the same task before being evaluated on the test samples. This evaluation setup tests the models’ ability to quickly adapt and learn from limited examples, simulating real-world scenarios where direct supervision may be scarce.

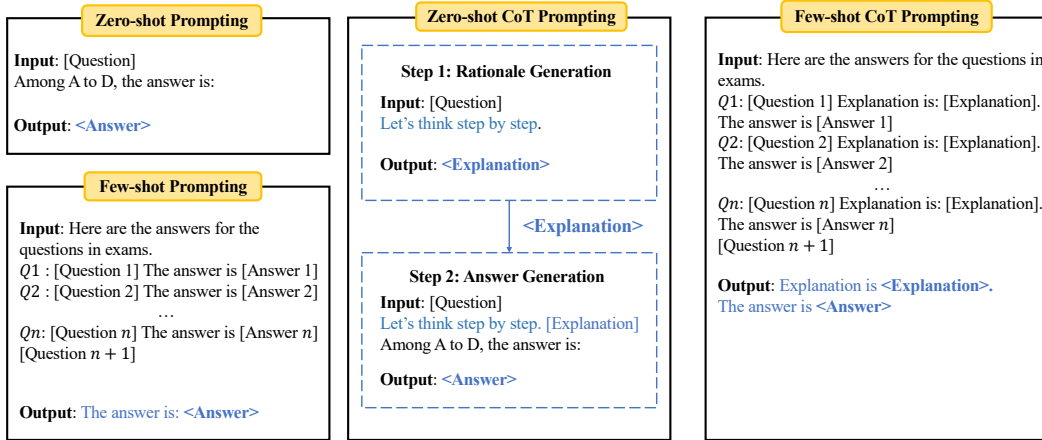


Figure 2: Prompting examples of different settings.

4.2.2 Chain-of-Thought (CoT) Reasoning

Chain-of-Thought (CoT) prompting (Wei et al. 2022) enables large language models to break down a complex question to a series of decomposed reasoning steps. To further probe the models’ reasoning capabilities, we implemented a “Chain-of-Thought” reasoning evaluation. This method involves asking the model to first generate an explanation for the given question, and subsequently, answer the question based on its self-generated explanation.

CoT는 2단계로 진행하며 아래와 같은 프롬프트로 구성

Step1
[Question]
순서대로 생각해봐라

Step2
[Question]
추가설명은: [Step1 Output]
답변은?

CoT 평가를 통해 모델의 Reasoning 능력을 평가함

More specifically, as shown in Fig. 2 the Chain-of-Thought reasoning evaluation is carried out in two steps for each task in the benchmark. In the first step, the model is prompted with “[question] Let’s think step by step: ” to generate a detailed explanation of its understanding of the problem and the reasoning process required to solve it (Zhang et al. 2022b). This allows us to evaluate the model’s ability to comprehend complex tasks and identify the essential elements needed for successful problem-solving. In the second step, the model is prompted with the generated explanation “[question] Explanation is: [explanation]. The answer is: ”, and provide an answer to the question, taking into account its self-generated explanation from the first step. This process aimed to test the

model's ability to use its own reasoning to derive a coherent and accurate solution, simulating the way humans often rely on their understanding and internal reasoning processes to solve problems. Specially, as shown in Fig. 2 in the few-shot CoT setting, the explanation and answer are generated in the same step.

4.2.3 Evaluation Metrics

To assess the performance of the language models on our benchmark, we used both quantitative and qualitative evaluation metrics. Quantitative metrics included task-specific accuracy for multi-choice questions and use Exact Match (EM) for fill-in-blank questions. For a more in-depth analysis, we also performed qualitative evaluations, which involved human evaluators assessing the models' responses in terms of semantic understanding capability, knowledge utilization, and reasoning quality. This combination of evaluation metrics allowed us to obtain a comprehensive understanding of the models' strengths and limitations in human-centric reasoning tasks.

Evaluation Metric은
양적, 질적 측정

양적 측정
- 다지선다형: Accuracy
- 빈칸채우기: Exact Match

질적 측정
사람이 직접 의미 이해 능력,
지식 활용, 추론 능력을 평가함

4.2.4 Implementation Details

All experiments were conducted using the respective language models' API provided by Azure OpenAI Service⁵. The Azure OpenAI services offer two types of APIs: completion and chat completion. The completion API generates text based on prompts, while the chat completion API generates the next AI response based on the conversation history and new human input. For Text-Davinci-003, we use the completion API, and for ChatGPT and GPT-4, we use the chat completion API. Notably, only the chat completion API is available for GPT-4 at present. We use a temperature of zero to generate output using greedy search and set the maximum number of tokens for generation to 2048. Additionally, we set the frequency penalty to zero and top p to 1, which are the default values for these APIs.

generate의 파라미터
- temperature = 0
- greedy search
- max_token = 2048
- frequency_penalty = 0
- top_p = 1

API의 디폴트 설정임

The Chat Completion API exhibits distinct properties in comparison to the Completion API. In a zero-shot context, the Chat Completion API has the potential to autonomously generate reasoning steps, eliminating the necessity for prompt engineering and potentially enhancing performance. For few-shot scenarios, it is imperative to adapt the few-shot examples into conversational history, as recommended in the Azure guidelines. The inquiry is transformed into a user input, while the AI's response is composed of a chain-of-thought explanation and answer. However, we have observed that the models, particularly ChatGPT, encounter difficulties in adhering to the pattern using the Chat Completion API. Consequently, we employ the Completion API to conduct few-shot experiments with ChatGPT, which is analogous to Text-Davinci-003, in order to gain a deeper understanding of the disparities between Text-Davinci-003 and ChatGPT. If a completion API for GPT-4 become accessible in the future, we will revise and update the few-shot outcomes accordingly.

긴요한

질의

차이

유사한

Few-shot Examples Construction: For AQuA-RAT, LogiQA and LSAT, we randomly sample five examples of medium sentence length of the test set from the provided training set. Similarly, for Gaokao and SAT, we randomly select five examples of medium sentence length from the dataset that was initially collected and exclude them from the test set. For JEC-QA, given that the test set is not publicly available, we take the first 1,000 examples from the training set as the test set and again sample five examples of medium sentence length from the rest. For MATH, we use the same instances as in the appendices of Lewkowycz et al. (2022).

AQuA-RAT, LogiQA, LSAT,
Gaokao, SAT의 경우
Training 셋에서 5개를 임의로
뽑아서 Few-shot의 example로
사용함

JEC-QA의 경우 테스트 셋이
없으므로,
Training Set에서 1000개의
샘플을 테스트셋으로 정하고
나머지 Training set에서
5개의 example을 샘플링

MATH도 똑같이함

To generate explanations for few-shot CoT experiments, for AQuA-RAT and MATH, we use the existing rationales from these datasets. For Gaokao and SAT, we collected expert annotations. For LogiQA, JEC-QA and LSAT, we use ChatGPT to generate explanations given the questions and the answers. We release all CoT demonstrations in the Github repository.

4.3 Main Results

The results of zero-shot learning and zero-shot CoT (Chain-of-Thought) are reported in Tab. 2. The results of few-shot learning and few-shot CoT are reported in Tab. 3. We also report average and top human performance on each task. From the results, we highlight the following findings.

⁵<https://azure.microsoft.com/en-us/products/cognitive-services/openai-service>

Table 2: Performance of LLMs on 20 tasks under **zero-shot** and **zero-shot CoT** settings. We also report human performance on each task. For LSAT, Gaokao and SAT, we report average (50%) and top (1%) human performance. The Text-Davinci-003 is abbreviated as TD-003.

Task/Model	Human Performance		Zero-Shot			Zero-Shot CoT		
	Avg.	Top	TD-003	ChatGPT	GPT-4	TD-003	ChatGPT	GPT-4
AQuA-RAT	85	100	29.9	31.9	40.6	42.1	55.9	73.2
MATH	40	90	11.9	26.4	35.7	19.1	31.9	47.7
LogiQA (English)	86	95	22.7	35.0	49.3	36.9	39.9	57.8
LogiQA (Chinese)	88	96	40.3	41.0	58.8	36.7	38.9	57.5
JEC-QA-KD	71	78	21.9	21.1	33.4	18.4	21.2	31.9
JEC-QA-CA	58	85	21.0	22.0	31.1	16.7	19.6	29.8
LSAT-AR	56	91	21.7	24.4	35.2	23.9	22.6	34.4
LSAT-LR	56	91	47.5	52.6	80.6	50.0	52.6	80.6
LSAT-RC	56	91	64.7	65.4	85.9	57.6	62.1	85.1
SAT-Math	66	94	35.5	42.7	64.6	54.6	70.9	95.0
SAT-English	66	94	74.8	81.1	88.8	75.7	77.7	85.9
SAT-English (w/o Psg.)	66	94	38.4	44.2	51.0	44.2	45.6	25.2
GK-Cn	65	85	43.9	39.0	53.3	35.4	33.7	44.7
GK-En	69	91	81.4	84.9	91.9	83.0	84.3	92.5
GK-geography	65	85	53.3	59.8	76.9	48.7	55.8	72.4
GK-history	64	85	47.3	59.7	77.4	37.0	50.2	76.5
GK-biology	68	89	40.5	52.9	75.7	30.0	42.4	71.9
GK-chemistry	66	86	27.1	38.7	51.7	24.6	33.8	52.2
GK-physics	71	94	22.0	33.0	39.0	18.5	29.5	45.5
GK-Math-QA	73	96	28.2	36.5	47.0	28.8	33.3	50.7
GK-Math-Cloze	73	96	17.0	7.6	16.1	4.2	5.1	15.3
Average	67	91	38.1	42.9	56.4	37.4	43.2	58.4

Table 3: Performance of LLMs on 20 tasks under **few-shot** and **few-shot CoT** settings. We also report human performance on each task. For LSAT, Gaokao and SAT, we report average (50%) and top (1%) human performance. The Text-Davinci-003 is abbreviated as TD-003.

Task/Model	Human Performance		Few-Shot			Few-Shot CoT		
	Avg.	Top	TD-003	ChatGPT	GPT-4	TD-003	ChatGPT	GPT-4
AQuA-RAT	85	100	30.3	31.1	50.8	47.2	60.6	74.0
MATH	40	90	10.3	14.8	15.1	15.1	30.1	25.3
LogiQA (English)	86	95	43.5	43.5	63.9	37.5	38.9	62.7
LogiQA (Chinese)	88	96	43.2	46.2	65.0	40.0	38.6	61.9
JEC-QA-KD	71	78	22.4	27.6	41.3	23.6	23.4	40.4
JEC-QA-CA	58	85	22.2	25.1	37.4	16.1	20.0	34.7
LSAT-AR	56	91	22.6	25.7	33.9	22.6	25.2	31.7
LSAT-LR	56	91	60.4	59.2	85.9	51.2	52.2	84.5
LSAT-RC	56	91	70.6	67.7	87.7	64.3	57.6	87.7
SAT-Math	66	94	44.6	40.9	71.8	55.5	65.0	89.6
SAT-English	66	94	84.0	81.1	88.8	76.7	78.2	85.9
SAT-English (w/o Psg.)	66	94	48.1	53.9	63.6	48.5	51.5	62.6
GK-Cn	65	85	25.6	41.5	61.4	29.3	37.8	51.6
GK-En	69	91	86.9	86.3	93.8	80.7	84.6	93.1
GK-geography	65	85	59.8	63.8	75.9	52.3	61.8	76.4
GK-history	64	85	49.0	57.6	77.8	51.9	58.4	78.2
GK-biology	68	89	44.3	52.4	80.0	32.9	50.0	72.9
GK-chemistry	66	86	32.4	44.0	54.6	35.8	33.8	54.1
GK-physics	71	94	31.0	33.5	43.5	27.5	36.5	54.5
GK-Math-QA	73	96	27.6	31.3	39.9	33.1	31.6	49.0
GK-Math-Cloze	73	96	5.9	5.9	11.0	5.93	8.5	16.1
Average	67	91	41.2	44.4	59.2	40.4	45	61.3

Table 4: Performance of open-source model Vicuna-13B under zero-shot and zero-shot CoT setting. Task names are abbreviated.

Task/Model			LogiQA		JEC-QA		LSAT			SAT			GK							
	AQA	MATH	En.	Cn.	KD	CA	AR	LR	RC	Math	En.	Cn	En	Geo.	His.	Bio.	Che.	Phy.	M.-QA	M.-Cloze
Vicuna (ZS)	26.4	6.8	18.4	23.5	14.3	12.4	22.2	25.5	30.5	24.6	50.5	25.6	50.7	24.6	28.9	20.5	26.6	15	22.5	2.5
Vicuna (ZS-CoT)	22.1	6.6	30.3	27.1	14.9	15.2	20.9	36.1	44.2	35.5	57.8	23.6	67	28.6	34.9	24.3	23.2	17	21.7	1.7

(1) **Superior Performance of GPT-4:** Our results indicate that, on average, GPT-4 significantly outperforms its counterparts (ChatGPT and Text-Davinci-003) across all four evaluation settings. Impressively, GPT-4 achieves 93.8% accuracy on Gaokao-English and 95% accuracy on SAT-MATH, demonstrating its superior general capabilities in handling human-centric tasks.

1) GPT4의 압도적인 성능

- Gaokao-English는 93.8%
- SAT-MATH 95%

Human-Centric Task에서
GPT4의 성능은 압도적임

(2) **Comparison between ChatGPT and Text-Davinci-003:** Our analysis shows that ChatGPT significantly outperforms Text-Davinci-003 in tasks that require a high degree of external knowledge, such as those involving geography, biology, chemistry, physics, and mathematics. This suggests that ChatGPT has a stronger knowledge base and is better equipped to handle tasks that necessitate a deep understanding of specific domains.

2) ChatGPT vs Davinci-003

ChatGPT가 특정 도메인의
지식에서 더 탁월함
(지리학, 생물학, 화학,
물리학, 수학 등)

On the other hand, ChatGPT slightly outperforms Text-Davinci-003 or achieves comparable results in tasks that require pure understanding and do not rely heavily on external knowledge, such as English and LSAT tasks, across all evaluation settings. This observation implies that both models are capable of handling tasks centered on language comprehension and logical reasoning without the need for specialized domain knowledge.

그러나 지식기반이 아닌 작업
(영어, LSAT) 등에서는 큰 차이가
없음
특정 도메인 지식이 필요 없는
언어 이해나 논리 추론은 둘다
가능해보임

(3) **Challenge of Complex Tasks:** Despite the overall good performance of the models, we observe that all the LLMs struggle with complex reasoning tasks, such as MATH, LSAT-AR, GK-physics, and GK-Math, across all evaluation settings. This highlights the limitations of these models in handling tasks that require advanced reasoning and problem-solving skills. The observed difficulties in tackling complex reasoning problems present an opportunity for future research and development, aiming to enhance the models' general reasoning capabilities.

3) Complex Task

모든 모델들이 MATH,
LSAT-AR, GK-Physics,
GK-Math 등
복잡한 Reasoning 작업을
어려워함

(4) **Few-shot Learning vs. Zero-shot Learning:** In our experiments, we observe that few-shot learning generally leads to only a limited improvement in performance compared to zero-shot learning. This finding suggests that the zero-shot capabilities of current large language models (LLMs) are approaching their few-shot learning abilities. This marks a significant advancement compared to the original GPT-3 (Brown et al., [2020]) model, where few-shot performance was considerably better than zero-shot. One plausible explanation for this development is the enhanced human-alignment and instruction tuning in current LLMs. These improvements enable the models to better understand the meaning and context of tasks in advance, thus allowing them to perform well even in zero-shot settings. This progress in LLMs' zero-shot capabilities highlights the effectiveness of recent advancements in instruction tuning of LLMs.

4) Few-shot vs Zero-shot

예전과 달리 Few-shot과
Zero-shot의 큰 차이가 없음

모델들이 Zero-shot 능력이
Few-shot에 근접함

또한 RLHF를 통해 모델들이
Task에 대한 이해를
잘하게 되었고 Zero-shot을
개선 시킨 것으로 보임

As shown in Fig. 4, Vicuna, despite excelling on OpenLLM leaderboard (Beeching et al., [2023]) and its claimed comparable ability with ChatGPT, falls short on AGIEval, highlighting the valuable challenges our benchmark presents to open-source models.

4.4 Analyses of Chain-of-thought Prompting (CoT)

As reported in Tab. 2 and Tab. 3, the Chain-of-Thought (CoT) prompting technique demonstrates its potential by improving average zero-shot and few-shot performance. However, the performance gains from CoT are not consistently observed across all tasks. Our analysis of CoT leads to the following findings:

CoT는 평균적으로
Zero-shot, Few-shot의
퍼포먼스를 올려줌
그러나 모든 Task에서
퍼포먼스를 높이는 건 아님

(1) **Performance Variability:** CoT substantially enhances performance in English mathematical exams, including MATH, AQuA-RAT, and SAT-Math. However, it leads to performance degradation in several other tasks, which may be a consequence of the model generating misleading reasoning processes. We also observed that the effects of CoT vary across different tasks, indicating that its impact on model performance is not uniform. It is crucial to further investigate the factors contributing to these performance variations and identify ways to optimize CoT for a broader range of tasks.

(2) **Backbone Dependency:** The effectiveness of CoT is influenced by the underlying model. For example, GPT-4 is better equipped to generate illustrative reasoning processes, which subsequently result in enhanced performance when employing CoT. This observation highlights the importance of considering the interplay between CoT and the backbone model, as their compatibility can significantly impact problem-solving capabilities.

2) Backbone Dependency
CoT의 퍼포먼스는 백본에 따라
달라짐 GPT4의 경우는 CoT의
효과가 컸음

(3) **Language Sensitivity:** The impact of CoT varies across different languages. In the case of the LogiQA exam, the English version is translated from the original Chinese version. Both ChatGPT and GPT-4 demonstrate performance improvements with CoT in LogiQA-English but suffer from performance degradation in LogiQA-Chinese. A similar pattern is observed in mathematical tests,

3) Language Sensitivity
CoT는 영어에서 더 잘함

LogiQA-Chinese에선
퍼포먼스가 떨어지는 반면,
LogiQA-English에선
퍼포먼스가 증가함

MATH, AQuA같은 영어로 된
수학 문제에서도 퍼포먼스 증가
반면 Gaokao에선 퍼포먼스 감소

where CoT boosts performance in English math exams (MATH, AQuA) but leads to performance declines in the Chinese math exam in Gaokao. These findings suggest that CoT's effectiveness is sensitive to language variations, emphasizing the need to further generalize and optimize CoT for different linguistic contexts. By tailoring CoT to better accommodate diverse languages, we can ensure more consistent and reliable problem-solving capabilities across a wider range of tasks and language settings.

These observations indicate that the effectiveness of CoT is relevant to task, model capability, and involved language. These factors should be carefully considered when employing CoT for specific tasks or developing future language models.

정리하자면,
CoT는 3가지 영향을 받음
1) Task에 따라
2) 모델의 능력에 따라
3) 언어에 따라

4.5 Qualitative Analyses of Model Capabilities

To gain a deeper understanding of the models' alignment with human capabilities, we perform a qualitative analysis by sampling 100 errorily answered instances for each task. Specifically, we examine the outputs from ChatGPT under the zero-shot CoT setting, where the model is required to generate both explanations and the corresponding answers. This approach allows us to assess the model's ability to handle the reasoning process and answer questions without further supervision.

ChatGPT에게 다음 요청
- Zero-Shot 수행
- 정답과 설명 요청

샘플로 100개의 에러를
질적분석 함
박사 과정 학생, 전문가 집단을
통해 모델을 4가지 측면에서
평가함

We enlist human annotators with expert knowledge, such as Ph.D. students and professional researchers, to evaluate the model outputs (i.e., explanations and answers) along the following four dimensions. The average scores on tasks for the four dimensions of capabilities are shown in Fig. 3.

Understanding
정확하게 Context와
Question에 대한 의미를
이해하고 있는지

Understanding: Evaluating whether the model can accurately comprehend the semantic meaning of the context and questions.

Knowledge
외부의 지식이나 공식을
정확히 생각할 수 있는지

Knowledge: Assessing whether the model can accurately recall relevant external knowledge or formulas for problem-solving.

Reasoning: Determining whether the model can perform correct reasoning steps or formulate accurate reasoning ideas to solve problems.

Reasoning
Reasoning Step과
Reasoning idea를 정확하게
공식화할 수 있는지

Calculation: Assessing whether the model can make correct mathematical calculations of the given formulas in the contexts of math, biology, chemistry and physics.

Calculation
수학, 생물학, 화학, 물리학을
풀기 위해 올바른 수학 계산을
할 수 있는지

An instance is scored with 1 if the model performs correctly on the corresponding skill and 0 otherwise. It is worth noting that some tasks, such as LSAT and certain English tasks, primarily emphasize understanding and reasoning without requiring extensive external knowledge or calculation. These tasks are excluded from the corresponding skill analyses to maintain focus on the relevant capabilities.

LSAT와 일부 English Task는
외부 지식이나 계산이 필요 없고
이해력과 추론만을 강조
=> 예제는 제외함

By conducting this qualitative analysis, we can obtain a more detailed understanding of the models' strengths and weaknesses in various aspects, shedding light on areas that may require further improvement in future iterations of large language models.

In addition to the capability scoring, we request annotators to provide a summary of the common detailed insights regarding the models' behavior patterns, emphasizing their strengths and weaknesses in addressing these human-centric tasks. We analyze the overall trends of the four aforementioned dimensions of capabilities in Section 4.5.1 outline the models' strengths in Section 4.5.2 and discuss their weaknesses in Section 4.5.3

4.5.1 Overall Trend of Model Capabilities

From the qualitative analysis in Fig. 3, we have the following observations:

Understanding: The model generally performs well in the understanding dimension. For most tasks, it can accurately interpret the meaning of questions, demonstrating its ability to comprehend context effectively.

모델들을 관찰해보니
다음 특징들이 나옴

Understanding
모델들이 대부분 문제와
문맥에 대한 이해는 잘함

Knowledge: In the knowledge dimension, the model demonstrates proficiency in identifying correct knowledge or formulas for mathematical and logical tasks. However, it encounters difficulties in recalling specific domain knowledge, such as law, biology, and physics. This observation emphasizes the significance of integrating more domain-specific knowledge into the model, potentially through the utilization of specialized domain-specific knowledge bases or knowledge-enhanced pre-training techniques.

Knowledge
특정 도메인에 대해 외부 지식을
생각해내는 것을 어려워함
(법, 생물학, 물리학 등)




Figure 3: Qualitative assessment of inaccurately answered questions by the model focuses on four dimensions of capabilities: understanding, knowledge acquisition, reasoning and calculation.

Reasoning: Among the four dimensions, the model’s reasoning capability appears to be relatively underdeveloped. For tasks necessitating complex, multi-step reasoning (e.g., LSAT-AR, LogiQA, and GK-Physics), the model struggles to accurately execute multi-step reasoning processes. This underlines the importance of future research concentrating on augmenting the model’s reasoning capabilities, potentially through the exploration of innovative prompting methods or training strategies that encourage complex reasoning and problem-solving skills.

Calculation: The model’s calculation ability is weaker than their understanding capacity and displays variability across different subjects. They perform better in math exams, but face challenges in chemistry and biology exams, which often require frequent variable substitution involving chemical elements. This suggests that enhancing the calculation and combinatorial abstraction and calculation ability of the model, particularly in subject areas with specialized notations or customized symbol substitutions, is a crucial challenge for further improvement.

Reasoning
대부분 Reasoning을 잘 못함
복잡하거나 여러 단계의 Reasoning을 필요로 하는 Task들 점수가 낮음
(LSAT-AR, LogiQA, GK-Physics 등)

Calculation
Task에 따라 다름
수학 시험은 잘하는데
화학과 생물학은 어려워 함

아마 화학 요소 같은 변수가 자주 바뀌기 때문인듯
전문적인 표기법을 기반으로 하는 과목에서 모델이 잘 조합하고 추상화하는 능력을 키우는게 좋아보임

These insights into the model's performance across the four dimensions provide valuable information to guide future research and development efforts aimed at enhancing large language models' capabilities in addressing human-centric tasks.

4.5.2 Strengths

By closely examining the models' output explanations and analyzing their behavior patterns, we identify several strengths that highlight the capabilities of these models in handling various aspects of problem-solving. The models demonstrate remarkable performance in the following areas:

Good Understanding: The models excel in ^{파악하다} accurately comprehending the semantic meaning of context and questions. They effectively ^{파악하다} discern nuances, interpret complex questions, and parse intricate sentences, showcasing their strong natural language understanding skills. This capability enables them to grasp the core concepts of a problem and lays the foundation for subsequent reasoning and problem-solving steps.

Proficiency in Simple Reasoning and Deduction: The models are ^{능숙한} adept at handling tasks that require simple reasoning and deduction. They can draw straightforward conclusions, identify logical connections, and perform basic inference, which is crucial for addressing a wide variety of problems. Their ability to effectively perform simple reasoning tasks is an essential component of their overall problem-solving skillset.

SOTA 모델들의 강점들
- 문제에 대한 이해를 잘함
- 간단한 추론은 잘함
- 일반적인 추론 과정(계획)을 잘 수립함

For example, the model can comprehend "could be true except" is equals to "cannot be true". Also, taking a question in the LSAT-AR task as an example, it requires the model to place 8 books to a bookcase with three shelves following conditions: "each shelf should have at least 2 books and more books should be placed on the bottom shelf than on the top shelf". The model successfully deduced that "there are at least 3 books on the bottom shelf and at most 2 books on the top shelf."

Grasping General Reasoning Process: The models demonstrate an ability to understand and generate the general idea of reasoning processes. They can identify the main components of a problem, recognize the structure of a solution, and outline a high-level reasoning strategy. This capability allows them to generate meaningful explanations and provides a starting point for more detailed reasoning and problem-solving tasks.

These strengths indicate that the models have made significant progress in aligning with human problem-solving capabilities. However, there is still room for improvement, especially in complex reasoning tasks and domain-specific knowledge, as discussed in the subsequent section on weaknesses.

4.5.3 Weaknesses

Despite the significant strengths displayed by the models, there are certain limitations that need to be addressed to improve their overall performance. We outline these weaknesses based on the analysis of the models' output explanations:

Understanding:

- Difficulty with Variable Substitution:** The models struggle to understand questions that require variable substitution, often failing to recognize the need for this operation and how it should be applied to solve the problem. This limitation can hinder their ability to tackle a wide range of mathematical and logical tasks. For instance, the model frequently struggles to answer chemistry questions that involve substituting a variable in a chemical equation with a chemical element and analyzing its properties.
- Challenges with Complex Math Concepts and Symbols:** The models find it difficult to comprehend complex mathematical concepts and interpret the meaning of symbols, particularly when multiple symbols are involved. This weakness limits their ability to effectively address advanced mathematical problems.
- Confusion with Similar Concepts:** The models can easily be confused by similar concepts or terms, sometimes leading to incorrect or misleading reasoning. For example, in the physics exam, the model is confused by the difference between vertical speed and horizontal speed of moving object. This issue underscores the need for better disambiguation and concept understanding techniques in future model iterations.

모델의 취약점들
Understanding
- 공식의 응용을 잘 못함
- 변수가 바뀌면 잘 못함
- 복잡한 수학 개념과 기호를 이해 못함
- 비슷한 개념을 혼동함
- Context가 길면 이해와 추론 능력이 떨어짐

- **Difficulty in Handling Long Contexts:** The models are prone to being disrupted by long contexts, leading to a decline in their comprehension and reasoning abilities. Improving the models' capacity to maintain focus and process extensive information is essential for enhancing their performance in real-world scenarios.

Knowledge:

- **Insufficiency in Commonsense and Domain-Specific Knowledge:** The models occasionally demonstrate a lack of commonsense or domain-specific knowledge, which **hinders** their ability to generate plausible explanations and provide accurate answers. This limitation underscores the importance of incorporating diverse knowledge sources into the training data and exploring techniques that can more effectively integrate and access this information within the models. Moreover, it emphasizes the necessity to broaden the models' exposure to a wider array of subjects and fields, ensuring a more comprehensive understanding of various domains.

For instance, given the conditions “if Julio and Kevin both lead morning sessions, we know that Kevin and Rebecca must lead sessions that meet on the same day,” the model incorrectly deduces that “Therefore, Rebecca must also lead a morning session.” This indicates a lack of commonsense knowledge about the relationship between *morning* and *day*, leading to an erroneous explanation. Additionally, the model generally performs poorly on tasks requiring specific domain knowledge, such as law and chemistry.

- **Difficulty Identifying Correct Formulas:** The models occasionally struggle to recall and apply the appropriate formulas necessary to solve particular problems, especially in tasks that demand specialized knowledge or expertise. This shortcoming suggests that there is potential for improvement in the models' knowledge retrieval mechanisms and their ability to recognize the relevance of specific formulas to a given problem. Developing strategies to enhance the models' proficiency in identifying and applying correct formulas will be essential for improving their performance in tasks requiring a deep understanding of domain-specific concepts and techniques.

Addressing these weaknesses in knowledge will contribute to the development of more robust and versatile large language models, better equipped to tackle a broader range of human-centric tasks and exhibit a more comprehensive understanding of various domains.

Reasoning:

- **Challenges in Strict Logical Deduction:** The models frequently encounter difficulties when attempting to perform strict logical deduction accurately. Common issues include ignoring premise conditions, misconstruing sufficient and necessary conditions, or making errors in logical chaining. These types of errors are commonly observed in manual analyses.

For instance, given a condition, “If Myers is on the team, neither Ortega nor Paine can be”, and a solution, “Ortega, Paine, Thomson, and Zayre are on the team”, the model incorrectly states that this solution is wrong because “Paine and Ortega are on the team”, neglecting to first satisfy the premise condition “If Myers is on the team”. Furthermore, the model demonstrates a misunderstanding of the difference between sufficient and necessary conditions in its explanation of another question and states: “If Kayne is assigned to an ambassadorship, then so is Jaramillo. This constraint is essentially the same as the given constraint that if Jaramillo is assigned to one of the ambassadorships, then so is Kayne”.

To address these limitations, it is essential to improve the models' abilities to recognize and apply logical rules and refine their understanding of logical structures.

- **Difficulty with Counterfactual Reasoning:** The models consistently struggle with counterfactual reasoning tasks. They have difficulty generating alternative scenarios, evaluating hypothetical outcomes, or exploring potential consequences based on varying assumptions. For instance, the models frequently make incorrect judgments for counterfactual questions in the LSAT-AR task: “Which one of the following, if substituted for the constraint that [Constraint A], would have the same effect in determining the assignment?” Enhancing the models' capabilities in handling counterfactual reasoning tasks is vital for developing a more comprehensive problem-solving skillset.
- **Struggles in Multi-hop Complex Reasoning:** The models have difficulty accurately executing multi-hop complex reasoning tasks, often displaying inconsistent logic, omitting inference

방해하다

Knowledge
- 가끔 Commonsense와 특정 도메인 지식의 부족을 보임
- 전문적 지식을 요구할 때 올바른 공식을 생각해내지 못함
(RAG로 해결하면 좋을듯)

sufficient condition : 충분조건
necessary condition : 필요조건

Reasoning
- 엄격한 논리적 추론을 잘 못함
전제조건을 무시하거나 충분조건과 필요조건을 혼동 논리적 연결에 오류를 범함

- 반사실적 추론의 어려움
대체 시나리오를 생성하거나 가상의 결과를 평가하거나 다양한 가정에 기반한 잠재적 결과를 탐구하는데 어려워 함

- 다중 단계의 복잡한 추론에서의 어려움
일관성 없는 논리를 보이거나 추론 단계를 생략하거나 잘못된 추론 체인을 생성함

steps, or producing flawed reasoning chains. To address a broader range of complex problems, it is crucial to improve the models' abilities to systematically navigate and process multi-step reasoning tasks.

- **Establishing Incorrect Conclusions and Contradictory Reasoning:** The models occasionally set an incorrect conclusion first and then generate contradictory reasoning based on that faulty foundation. This behavior emphasizes the need for improved reasoning verification and error correction techniques in the models' problem-solving processes.

- 잘못된 결론 수립과 모순된 추론
모델이 가끔 잘못된 결론을 설정하고 그 근거에 기반한 잘못된 추론을 생성함

- **Concealed Substitution of Concepts:** The models sometimes covertly substitute one concept with another similar one, leading to inaccurate or misleading reasoning. For example, in a biology exam, the model replaces the concept of "isotopically labeled amino acids" with "isotopically labeled tRNA (a tool for transporting amino acids)", resulting in erroneous reasoning. This issue underscores the importance of better concept disambiguation and reasoning coherence in future model iterations.

- 은밀히 개념을 교체함
가끔 모델들이 은밀하게 유사한 다른 컨셉으로 개념을 교체해버림

- **Difficulty in Identifying Solutions:** The models occasionally struggle to discover feasible solutions for specific problems, possibly due to limitations in their knowledge, reasoning capabilities, or problem-solving strategies. Addressing this shortcoming involves refining the models' ability to explore, evaluate, and select appropriate solutions based on the given problem context.

- 해결책 식별의 어려움
가끔 모델은 실현 가능한 해결책을 찾는데 어려워 함
모델의 지식, 추론 능력, 문제 해결 전략의 한계 때문인듯

- **Vulnerability to Contextual Disturbance:** The reasoning ability of large language models is often easily disrupted by changes in the surrounding context. When the context is modified, the models may produce different deductions for the same condition, suggesting that the robustness of their reasoning ability is not yet sufficient. This observation emphasizes the need to develop models that can maintain consistent reasoning performance, even in the presence of varying contextual information, ensuring more reliable and stable problem-solving capabilities.

- 문맥 변화에 대한 취약성
문맥이 변경되면 동일한 조건에서 다른 추론을 생성

Calculation: The model is prone to making calculation errors, particularly when dealing with complex variable substitutions. This may be attributed to the inherent limitations of the model's computation process in handling mathematical operations, as well as its difficulty in parsing intricate relationships between variables. Consequently, the model may struggle to maintain accuracy and precision when attempting to solve problems involving advanced algebraic manipulations or multi-step calculations. To address this issue, future iterations of the model should focus on enhancing its mathematical reasoning capabilities and improving its ability to recognize and apply relevant mathematical rules. This could involve incorporating specialized modules or mechanisms specifically designed to handle complex calculations, variable substitutions, and numerical problem-solving tasks. By refining the model's ability to accurately process and solve intricate mathematical problems, we can expand its applicability across a broader range of disciplines and domains, ensuring a more comprehensive and robust problem-solving skillset.

~의 결과로 보다

Calculation
특히 복잡한 변수들간의 교체를 다룰 때 오류를 범함
모델이 변수들간의 관계를 잘 이해하지 못해서 발생하는 것으로 보임

특히 대수학의 조작, 다단계 계산을 포함한 문제를 해결할 때 어려워함

By addressing these reasoning weaknesses, future large language models can be developed with more robust problem-solving capabilities, enabling them to effectively tackle a broader range of human-centric tasks and exhibit more sophisticated reasoning skills that align closely with human cognition.

5 Discussion about Future Directions

In light of the findings and limitations identified in our analysis, we point out several potential future directions for the development of large foundation models. These directions aim to address the weaknesses observed and further improve the models' capabilities in various human-centric tasks.

Inclusion of External Knowledge and Formulas: Enriching the models with external knowledge sources, like formulas and domain-specific knowledge can help enhance their performance in mathematical and knowledge-intensive tasks. Specifically, developing models that can effectively handle domain-specific tasks, such as those in law, biology, or physics, requires the integration of specialized knowledge bases and expertise into the model, and enables the model to adapt to different verticals more effectively. This could involve integrating structured knowledge repositories, mathematical and scientific concepts into the models with pre-training or knowledge-enhanced prompting methods, allowing them to access and apply relevant information more efficiently.

외부 지식 및 공식의 통합
법, 생물학, 물리학 등
도메인별 작업을 위해
Repository나
특정 지식 베이스를
모델과 통합시키는 게
좋은 방안이 될 수 있음

Strict Complex Logical Reasoning: Improving the models' capacity for strict complex logical reasoning is crucial for their performance in a wide range of human-centric tasks. This could involve the creation of new datasets that emphasize complex reasoning, as well as incorporating APIs and external symbolic compilers that can execute strict logical or mathematical deduction, and use the execution results to further facilitate logical analysis and reasoning verification.

엄격하고 복잡한 논리적 추론
외부 API나 새로운 데이터셋을
통한 학습이 방안이 될 수 있음

Multi-lingual Reasoning Capabilities Generalization: As mentioned in Sec. 4.4, the reasoning capabilities of models are variant across different language, where the reasoning ability is relatively better for rich-resourced language like English. Enhancing the models' multi-lingual reasoning capabilities is essential for their applicability in a diverse range of real-world scenarios. Therefore, future directions can put more focus on enhancing multilingual generalization of reasoning capability of foundation models.

다양한 언어를 기반으로 하는
Reasoning 능력

Multi-modal Evaluation: Expanding the evaluation framework to include multi-modal tasks can provide a more comprehensive assessment of the models' capabilities. This could involve incorporating visual, auditory, or interactive tasks that require the models to process and reason with multiple types of input simultaneously and generates multi-modal outputs for comprehensive real-world applications. In the future work, we will focus on the multi-modal version of AGIEval.

Multi Modal

Better Automatic Evaluation Metrics for Human-centric Tasks: Developing more robust and meaningful automatic evaluation metrics is crucial for the objective assessment of large language models' performance. Future research should focus on devising metrics that can accurately capture the models' understanding, knowledge, and reasoning abilities, while taking into account the nuances and complexities of real-world tasks.

자동화 Metric의 개발

Robustness of Reasoning Capability: Improving the robustness of the models' reasoning capabilities is essential for ensuring their consistency and reliability across various contexts. This can be achieved by exploring techniques that enhance the models' ability to maintain consistent reasoning performance, even when faced with changes in the surrounding context or variations in the input data.

Reasoning 과정을
견고하게 하도록 함
Input data가 바뀌더라도
모델의 Reasoning이
바뀌지 않도록

By addressing these future directions, foundation models can be further developed and refined to exhibit more advanced capabilities that align closely with human cognition, ultimately enabling them to tackle a broader range of complex, human-centric tasks with greater accuracy and reliability.

6 Conclusion

In this paper, we introduce AGIEval, a novel benchmark specifically designed to assess the general capabilities of large foundation models with respect to human-level cognition. The benchmark comprises high-quality official admission tests, qualification exams, and advanced competitions tailored for human participants, including law school admission tests and college entrance examinations. These assessments establish officially recognized standards for gauging human capabilities, making them well-suited for evaluating foundation models in the context of human-centric tasks. Additionally, AGIEval incorporates bilingual tasks in both Chinese and English, offering a more comprehensive assessment of model behavior. We have carried out an extensive evaluation of three cutting-edge large foundation models: Text-Davinci-003, ChatGPT, and GPT-4, using AGIEval. Remarkably, GPT-4 surpasses average human performance on LSAT, SAT, and math competition, attaining a 95% accuracy rate on the SAT Math test and a 92.5% accuracy on the Gaokao English test, demonstrating the impressive performance of contemporary foundation models. Despite their significant achievements, our in-depth manual analyses also reveal the limitations of these large language models in terms of understanding, knowledge utilization, reasoning and calculation. Guided by these findings, we explore potential future research avenues in this domain. By assessing these foundation models on human-centric tasks and probing their capabilities more deeply, we strive to foster the development of models that are more closely aligned with human cognition. Ultimately, this will enable them to tackle a broader range of intricate, human-centric tasks with increased accuracy and reliability.

AGIEval은
- 대학 입학시험
- 로스쿨 입학 시험
- 경시대회
등의 Human-Centric
문제로 구성됨

또한 영어 / 중국어의
Bilingual 벤치마크

GPT3.5, ChatGPT, GPT4로
평가했을 때,
GPT4는 LSAT, SAT,
수학 경시대회에서 평균적인
인간 수준을 넘어섬

또한 모델들은
- Knowledge Utilization
- Reasoning
- Calculation
에서 개선사항이 필요함

References

Edward Beeching, Sheon Han, Nathan Lambert, Nazneen Rajani, Omar Sanseviero, Lewis Tunstall, and Thomas Wolf. Open llm leaderboard. https://huggingface.co/spaces/HuggingFaceH4/open_llm_leaderboard, 2023.

- Samuel R. Bowman, Gabor Angeli, Christopher Potts, and Christopher D. Manning. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pp. 632–642, Lisbon, Portugal, September 2015. Association for Computational Linguistics. doi: 10.18653/v1/D15-1075. URL <https://aclanthology.org/D15-1075>
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrike, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*, 2023.
- Wei-Lin Chiang, Zhuohan Li, Zi Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E. Gonzalez, Ion Stoica, and Eric P. Xing. Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality, March 2023. URL <https://lmsys.org/blog/2023-03-30-vicuna/>.
- Jonathan H Choi, Kristin E Hickman, Amy Monahan, and Daniel Schwarcz. Chatgpt goes to law school. *Available at SSRN*, 2023.
- Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, et al. Scaling instruction-finetuned language models. *arXiv preprint arXiv:2210.11416*, 2022.
- Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. Boolq: Exploring the surprising difficulty of natural yes/no questions. *arXiv preprint arXiv:1905.10044*, 2019.
- Alexis Conneau and Douwe Kiela. SentEval: An evaluation toolkit for universal sentence representations. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 2018. European Language Resources Association (ELRA). URL <https://aclanthology.org/L18-1269>
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL <https://aclanthology.org/N19-1423>
- Jwala Dhamala, Tony Sun, Varun Kumar, Satyapriya Krishna, Yada Pruksachatkun, Kai-Wei Chang, and Rahul Gupta. Bold: Dataset and metrics for measuring biases in open-ended language generation. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pp. 862–872, 2021.
- Zhengxiao Du, Yujie Qian, Xiao Liu, Ming Ding, Jiezhong Qiu, Zhilin Yang, and Jie Tang. Glm: General language model pretraining with autoregressive blank infilling. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 320–335, 2022.
- Maud Ehrmann, Matteo Romanello, Simon Clematide, and Antoine Doucet. Introducing the HIPE 2022 Shared Task: Named Entity Recognition and Linking in Multilingual Historical Documents. In *Proceedings of the 44th European Conference on IR Research (ECIR 2022)*, Stavanger, Norway, 2022. Lecture Notes in Computer Science, Springer. URL https://link.springer.com/chapter/10.1007/978-3-030-99739-7_44
- Thomas Hartvigsen, Saadia Gabriel, Hamid Palangi, Maarten Sap, Dipankar Ray, and Ece Kamar. Toxigen: A large-scale machine-generated dataset for adversarial and implicit hate speech detection. *arXiv preprint arXiv:2203.09509*, 2022.

- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *Sort*, 2(4): 0–6.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020.
- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:452–466, 2019. doi: 10.1162/tacl_a_00276. URL <https://aclanthology.org/Q19-1026>
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving quantitative reasoning problems with language models, 2022.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022.
- Wang Ling, Dani Yogatama, Chris Dyer, and Phil Blunsom. Program induction by rationale generation: Learning to solve and explain algebraic word problems. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 158–167, 2017.
- Jian Liu, Leyang Cui, Hanmeng Liu, Dandan Huang, Yile Wang, and Yue Zhang. Logiqa: a challenge dataset for machine reading comprehension with logical reasoning. In *Proceedings of the Twenty-Ninth International Conference on International Joint Conferences on Artificial Intelligence*, pp. 3622–3628, 2021.
- Gary Marcus and Ernest Davis. *Rebooting AI: Building artificial intelligence we can trust*. Vintage, 2019.
- Bryan McCann, Nitish Shirish Keskar, Caiming Xiong, and Richard Socher. The natural language decathlon: Multitask learning as question answering. *arXiv preprint arXiv:1806.08730*, 2018.
- Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *Conference on Empirical Methods in Natural Language Processing*, 2018.
- OpenAI. Gpt-4 technical report, 2023.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35: 27730–27744, 2022.
- Denis Paperno, Germán Kruszewski, Angeliki Lazaridou, Quan Ngoc Pham, Raffaella Bernardi, Sandro Pezzelle, Marco Baroni, Gemma Boleda, and Raquel Fernández. The lambada dataset: Word prediction requiring a broad discourse context. *arXiv preprint arXiv:1606.06031*, 2016.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. Squad: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pp. 2383–2392, 2016.
- SenseTime. Internlm, 2023. <https://github.com/InternLM/InternLM-techreport/>.

- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. FEVER: a large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 809–819, New Orleans, Louisiana, June 2018. Association for Computational Linguistics. doi: 10.18653/v1/N18-1074. URL <https://aclanthology.org/N18-1074>
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pp. 353–355, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-5446. URL <https://aclanthology.org/W18-5446>
- Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. Superglue: A stickier benchmark for general-purpose language understanding systems. *Advances in neural information processing systems*, 32, 2019.
- Siyuan Wang, Zhongkun Liu, Wanjun Zhong, Ming Zhou, Zhongyu Wei, Zhumin Chen, and Nan Duan. From lsat: The progress and challenges of complex reasoning. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 30:2201–2216, 2022.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Ed Chi, Quoc Le, and Denny Zhou. Chain of thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022a.
- Zhuosheng Zhang, Aston Zhang, Mu Li, and Alex Smola. Automatic chain of thought prompting in large language models. *arXiv preprint arXiv:2210.03493*, 2022b.
- Haoxi Zhong, Chaojun Xiao, Cunchao Tu, Tianyang Zhang, Zhiyuan Liu, and Maosong Sun. Jec-qa: A legal-domain question answering dataset. In *Proceedings of AAAI*, 2020.
- Wanjun Zhong, Siyuan Wang, Duyu Tang, Zenan Xu, Daya Guo, Yining Chen, Jiahai Wang, Jian Yin, Ming Zhou, and Nan Duan. Analytical reasoning of text. In *Findings of the Association for Computational Linguistics: NAACL 2022*, pp. 2306–2319, Seattle, United States, July 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.findings-naacl.177. URL <https://aclanthology.org/2022.findings-naacl.177>

A Data Examples

Few data examples in Gaokao is shown in Fig. 4 and an example in SAT and corresponding Chain-of-Thought reasoning process generated by GPT-4 is shown in Fig. 5

B Data Contamination Issue

The issues surrounding data contamination and future web scrapes on training data for LLMs are noteworthy. Most of current benchmarks and datasets up to date suffer from these vulnerabilities. To exam the situation of contamination, we provided timestamp for the 9 new Gaokao datasets and we can evaluate on the latest tests (later than 2022) released later than the training data timestamp of ChatGPT and GPT-4. Hereinafter, from AGIEval, we provide results comparing the GPT-4 zero-shot performance on six Gaokao subjects with and without risk of data contamination (Chinese, English, and History have not been included in this analysis due to the constrained size of the exams for these subjects). The uncontaminated dataset comprises entries released in 2022, which postdates the

Example in Gaokao-MathQA

Question: 设 O 为坐标原点, 直线 $x = a$ 与双曲线 $C: \frac{x^2}{a^2} - \frac{y^2}{b^2} = 1 (a > 0, b > 0)$

的两条渐近线分别交于 D, E 两点, 若 $\triangle ODE$ 的面积为 8, 则 C 的焦距的最小值为 () ?

(Let O be the origin of the coordinate system, and let the line $x = a$ intersect the two asymptotes of the hyperbola $C: \frac{x^2}{a^2} - \frac{y^2}{b^2} = 1 (a > 0, b > 0)$ at points D and E . If the area of triangle $\triangle ODE$ is 8, what is the minimum value of the focal length of C ?)

Options: (A)4, (B)8, (C)16, (D)32

Answer: (B)

Example in Gaokao-Biology

Question: 人体下丘脑具有内分泌功能, 也是一些调节中枢的所在部位。下列有关下丘脑的叙述, 错误的是选项 (The hypothalamus in the human body has endocrine functions and is also the location of some regulatory centers. Which of the following statements about the hypothalamus is incorrect?)

Options:

(A) 下丘脑能感受细胞外液渗透压的变化 (The hypothalamus can sense changes in the osmotic pressure of extracellular fluid)

(B) 下丘脑能分泌抗利尿激素和促甲状腺激素 (The hypothalamus can secrete antidiuretic hormone and thyroid-stimulating hormone)

(C) 下丘脑参与水盐平衡的调节: 下丘脑有水平衡调节中枢 (The hypothalamus is involved in the regulation of water-salt balance; the hypothalamus has a water balance regulation center)

(D) 下丘脑能感受体温的变化: 下丘脑有体温调节中枢 (The hypothalamus can sense changes in body temperature; the hypothalamus has a body temperature regulation center)

Answer: (B)

Figure 4: Data examples in Gaokao.

Example in SAT-Math

Question: $3x + 4y = -23$
 $2y - x = -19$ What is the solution (x, y) to the system of equations above?

Options: (A). (-5,-2) (B). (3,-8) (C). (4,-6) (D). (9,-6)

Answer: (B)

Chain-of-Thought Reasoning by GPT-4

Step 1: Rewrite the second equation to isolate x .

Given: $3x + 4y = -23$
 $2y - x = -19$ Rewrite the second equation as $x = 2y + 19$

Step 2: Substitute the expression for x from the second equation into the first equation.

Substitute x in the first equation: $3(2y + 19) + 4y = -23$

Step 3: Solve for y .

Distribute the 3: $6y + 57 + 4y = -23$

Combine like terms: $10y + 57 = -23$

Subtract 57 from both sides: $10y = -80$

Divide by 10: $y = -8$

Step 4: Substitute the value of y back into the expression for x . $x = 2(-8) + 19$

Step 5: Solve for x . $x = 3$

Step 6: Write the solution as an ordered pair (x, y)

The solution to the system of equations is $(3, -8)$, which corresponds to answer choice (B).

Figure 5: Data example in SAT and corresponding CoT reasoning process.

GPT-4 training data's timestamp (September 2021). The results are reported on Tab. 5. Evidently, we observe that barring the Mathematics subjects, the performance experiences a minor drop in the absence of contamination, yet remains proximate to the performances on the complete datasets. This finding substantiates that while AGIEval may similarly encounter data contamination, it still retains its value as a useful and effective human-centric benchmark for evaluating the abilities of foundation models against complex human-oriented tasks. 가장 가까운

Data Contamination을 테스트해보기 위해
GPT4의 Knowledge Cutoff 이후의 데이터로 Reasoning을 함
(22p Table 5)

수학을 제외하고 퍼포먼스가 미세하게 떨어짐
그러나, 큰 차이는 없음
따라서 AGIEval의 DataContamination 이슈는 없어보임

Table 5: Analysis on data contamination risk on AGIEval. The uncontaminated set includes examples released later than the time stamp of training data of ChatGPT and GPT-4.

	full-set-size	accuracy	uncontaminated-set-size	accuracy
Gaokao-geography	199	76.9%	37	73%
Gaokao-biology	210	75.7%	58	77.6%
Gaokao-chemistry	207	51.7%	64	42.2%
Gaokao-physics	200	39%	5	40%
Gaokao-MathQA	351	47%	47	29.8%
Gaokao-mathcloze	118	16.1%	25	4%